

# cassava-leaf-disease-detection

January 31, 2024

## 0.1 Import Modules

```
[5]: import pandas as pd
import numpy as np
import time
import os
import copy
import json

# visualization modules
import seaborn as sns
import matplotlib.pyplot as plt
from PIL import Image

# pytorch modules
import torch
import torch.nn as nn
import torch.optim as optim
from torch.nn import functional as F
from torch.utils.data import Dataset, DataLoader
from torchvision import models
import torchvision.transforms as transforms

# augmentation
import albumentations
from albumentations.pytorch.transforms import ToTensorV2

import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

## 0.2 Load the Dataset

```
[6]: BASE_DIR = "../input/cassava-leaf-disease-classification/"

train = pd.read_csv(BASE_DIR+'train.csv')
train.head()
```

```
[6]:      image_id  label
0  1000015157.jpg      0
1  1000201771.jpg      3
2   100042118.jpg      1
3  1000723321.jpg      1
4  1000812911.jpg      3
```

```
[7]: # loading mapping for target label
with open(BASE_DIR+'label_num_to_disease_map.json') as f:
    mapping = json.loads(f.read())
    mapping = {int(k): v for k, v in mapping.items()}
mapping
```

```
[7]: {0: 'Cassava Bacterial Blight (CBB)',
      1: 'Cassava Brown Streak Disease (CBSD)',
      2: 'Cassava Green Mottle (CGM)',
      3: 'Cassava Mosaic Disease (CMD)',
      4: 'Healthy'}
```

```
[8]: train['label_names'] = train['label'].map(mapping)
train.head()
```

```
[8]:      image_id  label      label_names
0  1000015157.jpg      0  Cassava Bacterial Blight (CBB)
1  1000201771.jpg      3  Cassava Mosaic Disease (CMD)
2   100042118.jpg      1  Cassava Brown Streak Disease (CBSD)
3  1000723321.jpg      1  Cassava Brown Streak Disease (CBSD)
4  1000812911.jpg      3  Cassava Mosaic Disease (CMD)
```

### 0.3 Exploratory Data Analysis

```
[9]: def plot_images(class_id, label, total_images=6):
    # get image ids corresponding to the target class id
    plot_list = train[train['label']==class_id].
    ↪sample(total_images)['image_id'].tolist()

    labels = [label for i in range(total_images)]
    # dynamically set size for subplot
    size = int(np.sqrt(total_images))
    if size*size < total_images:
        size += 1

    # set figure size
    plt.figure(figsize=(15,15))

    # plot the image in subplot
    for index, (image_id, label) in enumerate(zip(plot_list, labels)):
```

```
plt.subplot(size, size, index+1)
image = Image.open(str(BASE_DIR+'train_images/'+image_id))
plt.imshow(image)
plt.title(label, fontsize=14)
plt.axis('off')

plt.show()
```

```
[10]: plot_images(0, mapping[0], 6)
```



```
[11]: plot_images(1, mapping[1], 6)
```

Cassava Brown Streak Disease (CBSD)



Cassava Brown Streak Disease (CBSD)



Cassava Brown Streak Disease (CBSD)



Cassava Brown Streak Disease (CBSD)



Cassava Brown Streak Disease (CBSD)



Cassava Brown Streak Disease (CBSD)



```
[12]: plot_images(2, mapping[2], 6)
```

Cassava Green Mottle (CGM)



Cassava Green Mottle (CGM)



Cassava Green Mottle (CGM)



Cassava Green Mottle (CGM)



Cassava Green Mottle (CGM)



Cassava Green Mottle (CGM)



```
[13]: plot_images(3, mapping[3], 6)
```



Cassava Mosaic Disease (CMD)



Cassava Mosaic Disease (CMD)



Cassava Mosaic Disease (CMD)



Cassava Mosaic Disease (CMD)



Cassava Mosaic Disease (CMD)



Cassava Mosaic Disease (CMD)



```
[14]: plot_images(4, mapping[4], 6)
```

Healthy



Healthy



Healthy



Healthy



Healthy




Healthy



```
[15]: # class distribution  
sns.countplot(train['label_names'])
```

```
plt.xticks(rotation=90)
```

```
[15]: (array([0, 1, 2, 3, 4]),  
      [Text(0, 0, 'Cassava Bacterial Blight (CBB)'),  
       Text(1, 0, 'Cassava Mosaic Disease (CMD)'),  
       Text(2, 0, 'Cassava Brown Streak Disease (CBSD)'),  
       Text(3, 0, 'Cassava Green Mottle (CGM)'),  
       Text(4, 0, 'Healthy')])
```



## 0.4 Configuration and Utility Functions

```
[16]: DIM = (256, 256)
      WIDTH, HEIGHT = DIM
      NUM_CLASSES = 5
      NUM_WORKERS = 24
      TRAIN_BATCH_SIZE = 32
      TEST_BATCH_SIZE = 32
      SEED = 1

      DEVICE = 'cuda'

      MEAN = [0.485, 0.456, 0.406]
      STD = [0.229, 0.224, 0.225]
```

## 0.5 Augmentations

```
[17]: def get_test_transforms(value = 'val'):
      if value == 'train':
          return augmentations.Compose([
              augmentations.Resize(WIDTH, HEIGHT),
              augmentations.HorizontalFlip(p=0.5),
              augmentations.Rotate(limit=(-90, 90)),
              augmentations.VerticalFlip(p=0.5),
              augmentations.Normalize(MEAN, STD, max_pixel_value=255.0,
          ↪always_apply=True),
              ToTensorV2(p=1.0)
          ])
      elif value == 'val':
          return augmentations.Compose([
              augmentations.Resize(WIDTH, HEIGHT),
              augmentations.Normalize(MEAN, STD, max_pixel_value=255.0,
          ↪always_apply=True),
              ToTensorV2(p=1.0)
          ])
```

## 0.6 Dataset Loader Class

```
[18]: class CassavaDataset(Dataset):
      def __init__(self, image_ids, labels, dimension=None, augmentations=None,
          ↪folder='train_images'):
          super().__init__()
          self.image_ids = image_ids
          self.labels = labels
          self.dim = dimension
          self.augmentations = augmentations
          self.folder = folder
```

```

# returns the length
def __len__(self):
    return len(self.image_ids)

# return the image and label for that index
def __getitem__(self, idx):
    img = Image.open(os.path.join(BASE_DIR, self.folder, self.
↪image_ids[idx]))

    if self.dim:
        img = img.resize(self.dim)

    # convert to numpy array
    img = np.array(img)

    if self.augmentations:
        augmented = self.augmentations(image=img)
        img = augmented['image']

    label = torch.tensor(self.labels[idx], dtype=torch.long)
    return img, label

```

## 0.7 Train Test Split

```

[19]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(train['image_id'],
↪train['label'], test_size=0.25)

```

```

[20]: from torch.utils.data import WeightedRandomSampler
def sampler_(labels):
    label_unique, counts = np.unique(labels, return_counts=True)
    print('Unique Labels', label_unique)
    weights = [sum(counts) / c for c in counts]
    sample_weights = [weights[w] for w in labels]
    sampler = WeightedRandomSampler(sample_weights, len(sample_weights),
↪replacement=True)
    return sampler

```

```

[21]: train_sampler = sampler_(y_train)

```

Unique Labels [0 1 2 3 4]

```

[22]: # create dataloaders for training antrain_test_splititidation
train_dataset = CassavaDataset(
    image_ids=x_train.values,
    labels=y_train.values,

```



```

        augmentations=get_test_transforms('train'),
        dimension=DIM
    )

    train_loader = DataLoader(
        train_dataset,
        batch_size=TRAIN_BATCH_SIZE,
        num_workers=NUM_WORKERS,
        shuffle=False,
        sampler=train_sampler
    )

    val_dataset = CassavaDataset(
        image_ids=x_test.values,
        labels=y_test.values,
        augmentations=get_test_transforms('val'),
        dimension=DIM
    )

    val_loader = DataLoader(
        val_dataset,
        batch_size=TRAIN_BATCH_SIZE,
        num_workers=NUM_WORKERS,
        shuffle=False
    )

    loaders = {'train': train_loader, 'val': val_loader}

```

```

[23]: # to check whether dataset is working or not
      # fetch the data based on index
      val_dataset[0]

```

```

[23]: (tensor([[[[ 0.5022,  0.8618,  0.4508, ..., -0.9192, -1.0219, -0.4054],
                 [-0.0458,  0.4337,  0.9303, ..., -0.7650, -0.6281, -0.6452],
                 [ 0.3652,  0.0398,  0.4337, ..., -1.4329, -1.1760, -1.0562],
                 ...,
                 [-0.0629, -0.2856, -0.2513, ..., -1.3987, -1.5185, -1.7069],
                 [-0.3712, -0.3541, -0.1999, ..., -1.4329, -1.5357, -1.6213],
                 [ 0.0056,  0.1254, -0.1657, ..., -1.3987, -1.5185, -1.6384]],
                [[ 0.5903,  0.9405,  0.5553, ..., -1.4055, -1.6506, -0.8803],
                 [-0.1275,  0.4853,  1.1856, ..., -1.0903, -1.0728, -0.9853],
                 [-0.2150, -0.4076,  0.2752, ..., -1.5630, -1.4755, -1.2654],
                 ...,
                 [ 0.8004,  0.4678,  0.3978, ..., -0.6176, -0.7227, -0.9153],
                 [ 0.3627,  0.3452,  0.4678, ..., -0.6702, -0.7577, -0.8277],
                 [ 0.5903,  0.7304,  0.4853, ..., -0.6527, -0.7402, -0.8452]]]])

```

```

[[[-0.1487,  0.3568, -0.0092, ..., -1.4210, -1.6302, -1.3164],
  [-0.7413, -0.0267,  0.6879, ..., -1.3164, -1.2293, -1.4907],
  [-0.5844, -0.6890,  0.0082, ..., -1.8044, -1.5953, -1.6476],
  ...,
  [-1.0027, -1.5081, -1.7870, ..., -1.4559, -1.4559, -1.5779],
  [-1.6127, -1.7173, -1.7347, ..., -1.3861, -1.4559, -1.5256],
  [-1.5953, -1.5430, -1.6999, ..., -1.2816, -1.4210, -1.5430]]]),
tensor(3))

```

## 0.8 Use Pretrained Model (Transfer Learning)

```

[24]: def getModel():
    net = models.resnet152(pretrained=True)

    # if you want to train the whole network, comment this code
    # freeze all the layers in the network
    for param in net.parameters():
        param.requires_grad = False

    num_fters = net.fc.in_features
    # create last few layers
    net.fc = nn.Sequential(
        nn.Linear(num_fters, 256),
        nn.ReLU(),
        nn.Dropout(0.3),
        nn.Linear(256, NUM_CLASSES),
        nn.LogSoftmax(dim=1)
    )

    # use gpu if any
    net = net.cuda() if DEVICE else net
    return net

```

```

[25]: model = getModel()

```

Downloading: "https://download.pytorch.org/models/resnet152-394f9c45.pth" to  
 /root/.cache/torch/hub/checkpoints/resnet152-394f9c45.pth

```

0%|          | 0.00/230M [00:00<?, ?B/s]

```

```

[26]: import math
def cyclical_lr(stepsize, min_lr=3e-4, max_lr=3e-3):

    # Scaler: we can adapt this if we do not want the triangular CLR
    scaler = lambda x: 1.

```

```

# Lambda function to calculate the LR
lr_lambda = lambda it: min_lr + (max_lr - min_lr) * relative(it, stepsize)

# Additional function to see where on the cycle we are
def relative(it, stepsize):
    cycle = math.floor(1 + it / (2 * stepsize))
    x = abs(it / stepsize - 2 * cycle + 1)
    return max(0, (1 - x)) * scaler(cycle)

return lr_lambda

```

```

[27]: criterion = nn.CrossEntropyLoss()
# optimizer = torch.optim.Adam(model.parameters(), lr=0.1)
optimizer = torch.optim.SGD(model.parameters(), lr=1., momentum=0.9)
step_size = 4*len(train_loader)
clr = cyclical_lr(step_size, min_lr=3e-4, max_lr=3e-3)
scheduler = torch.optim.lr_scheduler.LambdaLR(optimizer, [clr])

```

```

[28]: # print(model)

```

```

[29]: # freeze (or) unfreeze all the layers
unfreeze = True # to freeze, set it as False
for param in model.parameters():
    param.requires_grad = unfreeze

```

```

[30]: # find total parameters and trainable parameters
total_params = sum(p.numel() for p in model.parameters())
print(f'{total_params:,} total parameters')
trainable_params = sum(p.numel() for p in model.parameters() if p.requires_grad)
print(f'{trainable_params:,} training parameters')

```

```

58,669,637 total parameters
58,669,637 training parameters

```

## 0.9 Steps for Training and Validation

```

[31]: def train_model(model, dataloaders, criterion, optimizer, num_epochs=5,
    ↪ scheduler=scheduler):
    # set starting time
    start_time = time.time()

    val_acc_history = []

    best_model_wts = copy.deepcopy(model.state_dict())
    best_acc = 0.0

    for epoch in range(num_epochs):

```

```

print(f'Epoch {epoch}/{num_epochs-1}')
print('-'*15)

# each epoch have training and validation phase
for phase in ['train', 'val']:
    # set mode for model
    if phase == 'train':
        model.train() # set model to training mode
    else:
        model.eval() # set model to evaluate mode

    running_loss = 0.0
    running_corrects = 0
    fin_out = []

    # iterate over data
    for inputs, labels in dataloaders[phase]:
        # move data to corresponding hardware
        inputs = inputs.to(DEVICE)
        labels = labels.to(DEVICE)

        # reset (or) zero the parameter gradients
        optimizer.zero_grad()

        # training (or) validation process
        with torch.set_grad_enabled(phase=='train'):
            outputs = model(inputs)
            loss = criterion(outputs, labels)

            _, preds = torch.max(outputs, 1)

            # back propagation in the network
            if phase == 'train':
                loss.backward()
                optimizer.step()
                scheduler.step()

            running_loss += loss.item() * inputs.size(0)
            running_corrects += torch.sum(preds == labels.data)

    # calculate loss and accuracy for the epoch
    epoch_loss = running_loss / len(dataloaders[phase].dataset)
    epoch_acc = running_corrects.double() / len(dataloaders[phase].
dataset)

# print loss and acc for training & validation

```

```

        print('{} Loss: {:.4f} Acc: {:.4f}'.format(phase, epoch_loss,
↪epoch_acc))

        # update the best weights
        if phase == 'val' and epoch_acc > best_acc:
            best_acc = epoch_acc
            best_model_wts = copy.deepcopy(model.state_dict())
        if phase == 'val':
            val_acc_history.append(epoch_acc)

    print()
    end_time = time.time() - start_time

    print('Training completes in {:.0f}m {:.0f}s'.format(end_time // 60,
↪end_time % 60))
    print('Best Val Acc: {:.4f}'.format(best_acc))

    # load best model weights
    model.load_state_dict(best_model_wts)
    return model, val_acc_history

```

```

[32]: # train the model
model, accuracy = train_model(model=model, dataloaders=loaders,
↪criterion=criterion, optimizer=optimizer, num_epochs=5, scheduler=scheduler)

```

Epoch 0/4

-----

train Loss: 1.2207 Acc: 0.5154  
val Loss: 0.6961 Acc: 0.7460

Epoch 1/4

-----

train Loss: 0.7270 Acc: 0.7339  
val Loss: 0.6009 Acc: 0.7867

Epoch 2/4

-----

train Loss: 0.6226 Acc: 0.7768  
val Loss: 0.5246 Acc: 0.8185

Epoch 3/4

-----

train Loss: 0.5751 Acc: 0.7982  
val Loss: 0.6119 Acc: 0.7964

Epoch 4/4

-----



```
train Loss: 0.5003 Acc: 0.8212
val Loss: 0.4741 Acc: 0.8452
```

```
Training completes in 30m 54s
Best Val Acc: 0.8452
```

```
[ ]: # save the model and model weights
torch.save(model, '/kaggle/working/best_model.h5')
torch.save(model.state_dict(), '/kaggle/working/best_model_weights')
```

```
[ ]: # freeze (or) unfreeze all the layers
unfreeze = True # to freeze, set it as False
for param in model.parameters():
    param.requires_grad = unfreeze
```

```
[ ]: # # unfreeze selected layers
# layers = list(range(5,7))
# i = 0
# for layer in model.children():
#     if i in layers:
#         for param in layer.parameters():
#             param.requires_grad = True
#     i += 1
```

```
[ ]: # find total parameters and trainable parameters
total_params = sum(p.numel() for p in model.parameters())
print(f'{total_params:,} total parameters')
trainable_params = sum(p.numel() for p in model.parameters() if p.requires_grad)
print(f'{trainable_params:,} training parameters')
```

```
[ ]: # # gives the number of layers
# for i, layer in enumerate(model.children()):
#     print(i)
```

## 0.10 Testing the Model

```
[33]: # empty the cache from cuda device
torch.cuda.empty_cache()
```

```
[34]: def predict(model, dataloader, device):
    # set mode to eval
    model.eval()
    fin_out = []

    with torch.no_grad():
        for images, targets in dataloader:
            images = images.to(device)
```

```
        targets = targets.to(device)

        outputs = model(images)

        fin_out.append(F.softmax(outputs, dim=1).detach().cpu().numpy())

    return np.concatenate(fin_out)
```

```
[35]: # steps for model prediction
device = torch.device('cuda') # if you don't have gpu, set it as cpu
model.to(device)
pred = predict(model, val_loader, device)
pred = pred.argmax(axis=1)
```

```
[36]: y_test.values[:5]
```

```
[36]: array([3, 1, 3, 3, 2])
```

```
[37]: pred[:5]
```

```
[37]: array([3, 4, 3, 3, 2])
```